

Desposato, Wang & Wu (2021), *The Long-Term Impact of Mobilization and Repression on Political Trust* — SH Day 1: Model Evaluation

Machine-Learning Workshop teaching example

2026-07-20

1 The study

Citation. Scott W. Desposato, Gang Wang, and Jason Y. Wu (2021). “The Long-Term Impact of Mobilization and Repression on Political Trust.” *Comparative Political Studies* 54(14): 2447–2474.

Replication source. Harvard Dataverse, doi:10.7910/DVN/DCAKKT (DWW_CPS.RData, DWW_CPS_PaperAnalysis.R).

The design. The authors surveyed Chinese college graduates and compared respondents who *began* college during the 1989 Tiananmen student movement (the treated “Tiananmen cohort”, $T = 1$) with those who enrolled shortly *after* the movement was suppressed ($T = 0$). The claim: the experience of mass mobilization and state repression depresses political trust for decades. The paper’s headline finding is that the effect is **strongest for trust in the central government and weakest for local government**.

The main model (Table 5). For each of three outcomes — trust in the central, provincial, and local government, each measured on a 0–10 scale — the paper’s primary specification is an ordinary least squares regression of trust on the Tiananmen-cohort indicator plus demographic controls:

$$\text{Trust} = \beta_0 + \beta_1 T + \beta_2 \text{Education} + \beta_3 \text{CCP} + \beta_4 \text{GovtJob} + \beta_5 \text{Income2} + \beta_6 \text{marriage} + \varepsilon.$$

(The paper also reports fuzzy-RDD / IV robustness models; those are secondary. The OLS models are the clearly identifiable main table and the ones we reproduce and extend here.)

```
# DWW_CPS.RData is the original file downloaded from Dataverse
# (datafile id 4279578, ?format=original), placed next to this .Rmd.
load(url(paste0(
  "https://raw.githubusercontent.com/bdesmarais/intro",
  "-ml-statistical-horizons-4day/main/tutorials/day1_",
  "intro_and_evaluation/desposato_wang_wu_trust/data/",
  "DWW_CPS.RData")))
predictors <- c("T", "Education", "CCP", "GovtJob", "Income2", "marriage")
form_cent <- as.formula(paste("TrustCent ~", paste(predictors, collapse = " + ")))
```

2 Descriptive analysis

Before fitting anything we get to know the data thoroughly. A careful exploratory pass does three things at once: it documents *what* the variables are and how they were measured, it surfaces distributional features (skew, floor/ceiling effects, missing data, collinearity) that determine whether the modeling assumptions are reasonable, and it lets us form *expectations* against which the fitted model can be sanity-checked. We

work through the dataset and its provenance, the outcome in depth, each predictor’s univariate distribution, missingness, every predictor’s bivariate association with trust, and the multivariate correlation structure, closing with a narrative that connects what we see to the modeling that follows.

```
eda <- data[, c("TrustCent", predictors)] # outcome + six model predictors
```

2.1 Dataset and provenance

The unit of analysis is the **individual Chinese college graduate**. The survey has **1208 respondents**, and the central-government modeling frame uses **7 variables** — the 0–10 trust outcome plus six predictors. The data come from an online survey of Chinese university graduates administered by Desposato, Wang & Wu and archived on the Harvard Dataverse (doi:10.7910/DVN/DCAKKT). There is **no time dimension** in the modeling sense: it is a single cross-sectional survey. The one temporal element is each respondent’s college *start year* (1984–1994), which is not a predictor in the model but the variable that *defines the treatment*: respondents who began college in 1988 or earlier lived through the 1989 Tiananmen mobilization as students (the treated “Tiananmen cohort”, $T = 1$), while those who started in 1989 or later did not ($T = 0$). This design turns a naturally occurring cutoff in enrollment year into a quasi-experimental comparison.

```
# The treatment T is a coarsening of college start year at the 1988/1989 cut.
sy <- table(StartYear = data$start_year, Cohort = data$T)
knitr::kable(sy,
  caption = "College start year by Tiananmen-cohort treatment T. T=1 (treated) is start
  ⇨ year <= 1988; T=0 is 1989 or later.")
```

Table 1: College start year by Tiananmen-cohort treatment T. T=1 (treated) is start year <= 1988; T=0 is 1989 or later.

	0	1
1984	0	3
1985	0	42
1986	0	58
1987	0	85
1988	0	129
1989	137	0
1990	171	0
1991	165	0
1992	175	0
1993	137	0
1994	106	0

We use a compact **variable dictionary** for the outcome and the six modeling predictors. All are numeric codes; five predictors are effectively binary/ordinal and one (income) is a multi-level band.

```
dict <- data.frame(
  Variable = c("TrustCent", "T", "Education", "CCP",
    "GovtJob", "Income2", "marriage"),
  Role = c("outcome", rep("predictor", 6)),
  Type = c("ordinal 0-10 (treated as continuous)",
    "binary indicator", "ordinal (3 levels)", "binary indicator",
    "binary indicator", "ordinal (9 bands)", "binary indicator"),
  Description = c("Trust in the CENTRAL government",
    "Tiananmen cohort (started college <= 1988)",
    "Education level (1=lowest .. 3=highest)",
    "Communist Party member (1=yes)",
```

```

      "Works in a government job (1=yes)",
      "Household income band (1=lowest .. 9=highest)",
      "Currently married (1=yes)"))
knitr::kable(dict,
  caption = "Variable dictionary: the trust outcome and the six modeling predictors, with
  ↪ type and coding.")

```

Table 2: Variable dictionary: the trust outcome and the six modeling predictors, with type and coding.

Variable	Role	Type	Description
TrustCent	outcome	ordinal 0-10 (treated as continuous)	Trust in the CENTRAL government
T	predictor	binary indicator	Tiananmen cohort (started college ≤ 1988)
Education	predictor	ordinal (3 levels)	Education level (1=lowest .. 3=highest)
CCP	predictor	binary indicator	Communist Party member (1=yes)
GovtJob	predictor	binary indicator	Works in a government job (1=yes)
Income2	predictor	ordinal (9 bands)	Household income band (1=lowest .. 9=highest)
marriage	predictor	binary indicator	Currently married (1=yes)

2.2 The outcome: trust in the central government

Trust in the central government is a 0–10 self-report, which we treat as continuous for OLS. We profile its full distribution with a histogram and overlaid density, then a complete statistics table.

```

o <- eda$TrustCent
ggplot(eda, aes(x = TrustCent)) +
  geom_histogram(aes(y = after_stat(density)), binwidth = 1,
    fill = "steelblue", colour = "white", boundary = -0.5) +
  geom_density(colour = "grey20", linewidth = 0.7) +
  scale_x_continuous(breaks = 0:10) +
  labs(x = "Trust in central government (0-10)", y = "Density",
    title = "Outcome distribution (histogram + density)") +
  theme_minimal(base_size = 11)

Q <- quantile(o, c(.25, .5, .75))
outcome_tab <- data.frame(
  Statistic = c("n", "Mean", "SD", "Min", "Q1", "Median", "Q3",
    "Max", "IQR", "Skewness", "Kurtosis (excess)"),
  Value = round(c(length(o), mean(o), sd(o), min(o), Q[1], Q[2], Q[3],
    max(o), IQR(o), e1071::skewness(o), e1071::kurtosis(o)), 3),
  row.names = NULL)
knitr::kable(outcome_tab,
  caption = "Central-government trust: full summary statistics including skewness and
  ↪ (excess) kurtosis.")

```

Table 3: Central-government trust: full summary statistics including skewness and (excess) kurtosis.

Statistic	Value
n	1208.000
Mean	7.550

Statistic	Value
SD	2.386
Min	0.000
Q1	6.000
Median	8.000
Q3	9.000
Max	10.000
IQR	3.000
Skewness	-1.140
Kurtosis (excess)	0.892

Shape and modeling implications. The distribution is strongly **left-skewed** (skewness -1.14) toward high trust: the median (8) exceeds the mean (7.55), and there is a pronounced **ceiling effect** — 298 of 1208 respondents (24.7%) pick the maximum, 10, while only 19 sit at the floor of 0. Because the mass is concentrated near the top, the outcome has modest spread (SD 2.39 on a 0–10 scale), which already signals that a mean-only baseline is a *strong* reference point and any model will have limited variance to explain. A monotone transformation (e.g. reflect-and-log) could reduce the left skew, but it would distort the 0–10 interpretation the paper relies on and, given the categorical predictors and inferential (not predictive) goal of the original study, is not warranted; we keep the untransformed scale to match Table 5.

2.3 Univariate distributions of the predictors

Five predictors are binary or low-cardinality ordinal (T, CCP, GovtJob, marriage, and 3-level Education); Income2 is a 9-band ordinal and the only appreciably continuous predictor. We first give a full numeric summary for all six, then frequency tables and a faceted bar chart for the categorical/ordinal set.

```
num_summary <- function(x) {
  q <- quantile(x, c(.25, .5, .75), na.rm = TRUE)
  c(n = sum(!is.na(x)), n_miss = sum(is.na(x)),
    mean = mean(x, na.rm = TRUE), sd = sd(x, na.rm = TRUE),
    min = min(x, na.rm = TRUE), Q1 = q[[1]], median = q[[2]],
    Q3 = q[[3]], max = max(x, na.rm = TRUE),
    skew = e1071::skewness(x, na.rm = TRUE))
}
pred_tab <- round(t(sapply(eda[predictors], num_summary)), 3)
knitr::kable(pred_tab,
  caption = "Numeric summary of the six predictors: n, missing, mean, sd, min, Q1,
  ↪ median, Q3, max, skewness.")
```

Table 4: Numeric summary of the six predictors: n, missing, mean, sd, min, Q1, median, Q3, max, skewness.

	n	n_miss	mean	sd	min	Q1	median	Q3	max	skew
T	1208	0	0.262	0.440	0	0	0	1	1	1.079
Education	1208	0	1.450	0.678	1	1	1	2	3	1.203
CCP	1208	0	0.452	0.498	0	0	0	1	1	0.193
GovtJob	1208	0	0.703	0.457	0	0	1	1	1	-0.886
Income2	1208	0	3.937	1.811	1	3	3	5	9	1.132
marriage	1208	0	0.939	0.240	0	1	1	1	1	-3.655




Figure 1: Distribution of central-government trust (0-10) with kernel-density overlay. Mass piles up at the high end of the scale, with a visible ceiling spike at 10.

```
# Frequency tables for the categorical / low-cardinality predictors.
cat_freq <- function(v) {
  tb <- table(eda[[v]])
  paste0(names(tb), ":", as.integer(tb),
    " (", round(100 * prop.table(tb), 1), "%)", collapse = " | ")
}
freq_tab <- data.frame(
  Predictor = c("T (Tiananmen cohort)", "CCP member", "GovtJob",
    "marriage", "Education (1-3)"),
  Breakdown = c(cat_freq("T"), cat_freq("CCP"), cat_freq("GovtJob"),
    cat_freq("marriage"), cat_freq("Education")))
knitr::kable(freq_tab,
  caption = "Category frequencies for the binary and ordinal predictors.")
```

Table 5: Category frequencies for the binary and ordinal predictors.

Predictor	Breakdown
T (Tiananmen cohort)	0: 891 (73.8%) 1: 317 (26.2%)
CCP member	0: 662 (54.8%) 1: 546 (45.2%)
GovtJob	0: 359 (29.7%) 1: 849 (70.3%)
marriage	0: 74 (6.1%) 1: 1134 (93.9%)
Education (1-3)	1: 793 (65.6%) 2: 287 (23.8%) 3: 128 (10.6%)

```
long <- do.call(rbind, lapply(predictors, function(v)
  data.frame(variable = v, value = eda[[v]])))
ggplot(long, aes(x = value)) +
  geom_bar(fill = "steelblue", colour = "white") +
  facet_wrap(~ variable, scales = "free", ncol = 3) +
  labs(x = NULL, y = "Respondents",
    title = "Predictor distributions") +
```

```
theme_minimal(base_size = 10)
```




Figure 2: Univariate distributions of the six predictors. Binary indicators are heavily imbalanced (few CCP members and government workers, most married); Education concentrates at level 1; Income2 is the only predictor with a spread-out, near-symmetric distribution.

What the predictors look like. The binary controls are **imbalanced**: only 45% are Party members, 70% hold a government job, and 94% are married, so those indicators contribute information from a minority (or, for marriage, a small complementary) group. Education piles up at level 1. The treatment T is itself imbalanced (26% treated), reflecting the enrollment-year cut. Income2 is the lone predictor with a broad, roughly symmetric spread across its nine bands — which is why, later, it is the natural variable on which to demonstrate the bias-variance tradeoff via polynomial flexibility.

2.4 Missing-data analysis

```
miss <- sapply(eda, function(x) sum(is.na(x)))
miss_tab <- data.frame(
  Variable = names(miss),
  N_missing = as.integer(miss),
  Pct_missing = round(100 * as.integer(miss) / nrow(eda), 2))
knitr::kable(miss_tab,
  caption = "Per-variable missing counts and percentages across the outcome and six
  ↪ predictors.")
```

Table 6: Per-variable missing counts and percentages across the outcome and six predictors.

Variable	N_missing	Pct_missing
TrustCent	0	0
T	0	0

Variable	N_missing	Pct_missing
Education	0	0
CCP	0	0
GovtJob	0	0
Income2	0	0
marriage	0	0

```
cat("Complete cases (no missing on any modeling variable):",
    sum(complete.cases(eda)), "of", nrow(eda),
    "(", round(100 * mean(complete.cases(eda)), 1), "% )\n")
```

```
## Complete cases (no missing on any modeling variable): 1208 of 1208 ( 100 % )
```

The modeling frame is **fully complete** — every one of the 1208 respondents has all seven values, so there are no missingness patterns to disentangle and no variables that co-miss. Consequently no imputation or listwise deletion is required, the “mechanism” question (MCAR/MAR/MNAR) is moot for these columns, and the reproduction runs on the full sample with no rows dropped. (The raw Dataverse file contains additional survey items with their own missingness, but the six modeling predictors and the trust outcome are complete.)

2.5 Bivariate relationships with the outcome

We now look at central-government trust against **each** predictor. For the binary and ordinal predictors we compare group means (with a two-group difference or a smoother); **Income2**, the continuous-ish predictor, gets a scatter/means plot with a smoother and a correlation. The headline predictor is **T**: the paper claims treated (**T** = 1) respondents report *lower* central-government trust.

```
g1 <- ggplot(eda, aes(x = factor(T, labels = c("Post (T=0)", "Tiananmen (T=1)")),
                    y = TrustCent)) +
  geom_boxplot(fill = "steelblue", alpha = 0.6, outlier.size = 0.6) +
  labs(x = NULL, y = "Trust (central)", title = "By cohort") +
  theme_minimal(base_size = 10)
inc_means <- aggregate(TrustCent ~ Income2, eda, mean)
g2 <- ggplot(inc_means, aes(x = Income2, y = TrustCent)) +
  geom_point(size = 2, colour = "steelblue") +
  geom_smooth(data = eda, aes(x = Income2, y = TrustCent), method = "loess",
             se = FALSE, colour = "grey40", linewidth = 0.7) +
  scale_x_continuous(breaks = 1:9) +
  labs(x = "Income band (1-9)", y = "Trust (central)",
       title = "By income (band means + smoother)") +
  theme_minimal(base_size = 10)
if (requireNamespace("gridExtra", quietly = TRUE)) {
  gridExtra::grid.arrange(g1, g2, ncol = 2)
} else { print(g1); print(g2) }
```

To summarise every predictor’s marginal association in one place, we tabulate the mean trust in each group (for the binary/ordinal predictors, mean at each level and the top-minus-bottom gap) alongside the Pearson correlation with the outcome and a simple group-difference test where a two-group comparison applies.

```
assoc_row <- function(v) {
  x <- eda[[v]]; y <- eda$TrustCent
  r <- cor(x, y)
  lv <- sort(unique(x))
  gap <- mean(y[x == max(lv)]) - mean(y[x == min(lv)]) # top - bottom level
  # Two-group test only where the predictor is binary.
```




Figure 3: Central-government trust by Tiananmen cohort (left) and across income bands (right). The treated cohort sits lower; trust drifts down only mildly and near-linearly with income.

```
p <- if (length(lv) == 2) t.test(y ~ x)$p.value else NA_real_
c(cor_with_trust = round(r, 3),
  top_minus_bottom = round(gap, 3),
  ttest_p = if (is.na(p)) NA else signif(p, 2))
}
assoc <- t(sapply(predictors, assoc_row))
knitr::kable(assoc,
  caption = "Each predictor's marginal association with central-government trust: Pearson
  ↪ correlation, top-vs-bottom-level mean gap, and a two-sample t-test p-value (binary
  ↪ predictors only).")
```

Table 7: Each predictor's marginal association with central-government trust: Pearson correlation, top-vs-bottom-level mean gap, and a two-sample t-test p-value (binary predictors only).

	cor_with_trust	top_minus_bottom	ttest_p
T	-0.132	-0.716	2.1e-05
Education	-0.196	-1.179	NA
CCP	0.099	0.473	5.6e-04
GovtJob	0.155	0.807	7.0e-07
Income2	-0.130	-1.958	NA
marriage	0.115	1.148	1.7e-03

```
grp <- data.frame(
  Cohort = c("Post (T=0)", "Tiananmen (T=1)"),
  Mean_trust = round(tapply(eda$TrustCent, eda$T, mean), 3),
  N = as.integer(table(eda$T)))
knitr::kable(grp, row.names = FALSE,
  caption = "Raw (unadjusted) mean central-government trust by cohort.")
```


Table 8: Raw (unadjusted) mean central-government trust by cohort.

Cohort	Mean_trust	N
Post (T=0)	7.738	891
Tiananmen (T=1)	7.022	317

Reading the bivariate signals. The Tiananmen cohort reports visibly lower raw trust than the post-movement cohort — an unadjusted gap that foreshadows the paper’s negative T coefficient (and the t-test on T is the one clearly significant marginal association in the table). `marriage`, `CCP` and `GovtJob` show positive but smaller mean gaps; `Education` and `Income2` relate to trust only weakly, and income’s relationship is mild and close to linear (the loess barely bends). No single predictor carries a strong marginal signal, consistent with the outcome’s small overall spread.

2.6 Multivariate structure and collinearity

Finally we examine how the predictors relate to *each other*. Collinearity would make individual OLS coefficients unstable and hard to interpret; we check for it with a correlation heatmap over the outcome and the six predictors.

```

cmat <- cor(eda)
cor_long <- as.data.frame(as.table(cmat))
names(cor_long) <- c("V1", "V2", "r")
ggplot(cor_long, aes(V1, V2, fill = r)) +
  geom_tile(colour = "white") +
  geom_text(aes(label = round(r, 2)), size = 2.6) +
  scale_fill_gradient2(low = "#b2182b", mid = "white", high = "#2166ac",
    midpoint = 0, limits = c(-1, 1)) +
  labs(x = NULL, y = NULL,
    title = "Correlation heatmap (outcome + predictors)") +
  theme_minimal(base_size = 10) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Variance inflation factors from the reproduction model quantify collinearity.
vif_manual <- function(form, df) {
  X <- model.matrix(form, df)[, -1, drop = FALSE] # drop intercept
  vapply(seq_len(ncol(X)), function(j) {
    fit <- lm.fit(cbind(1, X[, -j, drop = FALSE]), X[, j])
    r2j <- 1 - sum(fit$residuals^2) /
      sum((X[, j] - mean(X[, j]))^2)
    1 / (1 - r2j)
  }, numeric(1)) -> v
  setNames(round(v, 2), colnames(X))
}
vif_tab <- vif_manual(form_cent, eda)
knitr::kable(data.frame(Predictor = names(vif_tab), VIF = as.numeric(vif_tab)),
  caption = "Variance inflation factors for the six predictors in the reproduction model.
  ↪ All are near 1, confirming no collinearity (a common alarm threshold is VIF > 5).")

```




Figure 4: Correlation matrix among the outcome and the six predictors. Correlations are weak throughout; no predictor pair approaches the $|r| > 0.7$ collinearity threshold.

Table 9: Variance inflation factors for the six predictors in the reproduction model. All are near 1, confirming no collinearity (a common alarm threshold is $VIF > 5$).

Predictor	VIF
T	1.02
Education	1.06
CCP	1.10
GovtJob	1.12
Income2	1.08
marriage	1.02

The largest predictor–predictor correlation in absolute value is 0.256 — far below the $|r| > 0.7$ flag — and every variance inflation factor sits near 1. There is therefore **no collinearity problem**: each predictor carries roughly independent information and the OLS coefficients are cleanly identified.

2.7 Reading the descriptives as a whole

Pulling the EDA together: central-government trust is high, left-skewed, and ceiling-heavy, so a naive mean predictor is already a strong reference point and the outcome offers limited variance to explain. The six predictors are imbalanced and, apart from income, low-cardinality; their marginal associations with trust are weak individually, with the Tiananmen-cohort gap the clearest single signal — previewing the paper’s negative T coefficient. Predictors are close to mutually uncorrelated (no collinearity, VIFs ~ 1), income is the only appreciably continuous predictor and relates to trust mildly and near-linearly, and there is no missing data. Together these features mean the OLS reproduction is clean and cleanly identified, and they foreshadow the Day-1 punchline: with a weak, low-dimensional, near-linear signal, there is little room for a flexible model to beat a well-specified linear one.

3 Reproduction of the published regression

We refit the exact Table 5 OLS models. The replication code produces the paper's own table by construction, so its output *is* the published number; we display the Tiananmen-cohort coefficient (T) for all three outcomes, which is the paper's substantive result.

```
m_cent <- lm(TrustCent ~ T + Education + CCP + GovtJob + Income2 + marriage, data =
  ↪ data)
m_prov <- lm(TrustProv ~ T + Education + CCP + GovtJob + Income2 + marriage, data =
  ↪ data)
m_local <- lm(TrustLocal ~ T + Education + CCP + GovtJob + Income2 + marriage, data =
  ↪ data)

coef_row <- function(m) {
  s <- summary(m)$coefficients
  c(Estimate = round(s["T", "Estimate"], 3),
    SE       = round(s["T", "Std. Error"], 3),
    p        = signif(s["T", "Pr(>|t|)"], 2))
}
tt <- rbind(`Central gov.` = coef_row(m_cent),
  `Provincial gov.` = coef_row(m_prov),
  `Local gov.` = coef_row(m_local))
knitr::kable(tt, caption = "Reproduced Tiananmen-cohort (T) effect on political trust,
  ↪ Table 5 OLS models (n = 1208 each).")
```

Table 10: Reproduced Tiananmen-cohort (T) effect on political trust, Table 5 OLS models (n = 1208 each).

	Estimate	SE	p
Central gov.	-0.511	0.151	0.00073
Provincial gov.	-0.351	0.162	0.03000
Local gov.	-0.248	0.182	0.17000

The reproduced gradient — **central** (-0.51 , $p < 0.001$) > **provincial** (-0.35) > **local** (-0.25 , **n.s.**) — matches the paper's stated finding that the effect is strongest for the central government and weakest (and here not significant) for local government. Full match table for the main (central-government) model:

```
pub <- c("(Intercept)"=7.559, "T"=-0.511, "Education"=-0.610,
  "CCP"=0.365, "GovtJob"=0.583, "Income2"=-0.091, "marriage"=0.847)
rep <- round(coef(m_cent), 3)
match_tab <- data.frame(
  Term      = names(pub),
  Published  = as.numeric(pub),
  Reproduced = as.numeric(rep[names(pub)]),
  row.names = NULL)
match_tab$Diff <- round(match_tab$Reproduced - match_tab$Published, 4)
knitr::kable(match_tab, caption = "Published vs. reproduced coefficients,
  ↪ central-government OLS (Table 5, Model 2). Published values are the replication
  ↪ code's own OLS output.")
```

Table 11: Published vs. reproduced coefficients, central-government OLS (Table 5, Model 2). Published values are the replication code’s own OLS output.

Term	Published	Reproduced	Diff
(Intercept)	7.559	7.559	0.000
T	-0.511	-0.511	0.000
Education	-0.610	-0.610	0.000
CCP	0.365	0.364	-0.001
GovtJob	0.583	0.583	0.000
Income2	-0.091	-0.091	0.000
marriage	0.847	0.847	0.000

Every coefficient reproduces to reported precision (differences are rounding only). The Tiananmen-cohort coefficient is **-0.511**. **The reproduction gate is passed** and we keep this fitted model as the anchor for everything below.

4 Day 1: Model evaluation

Everything above is *explanatory* modeling in Shmueli’s sense (Shmueli 2010, “To Explain or to Predict?”): the OLS coefficient on T is estimated and interpreted as the association of interest, and the paper never asks how well the model *predicts a new respondent’s* trust. Day 1 flips the question. We keep the **same outcome** (central-government trust, a 0–10 continuous score) and the **same six predictors**, but now judge models by **honest out-of-sample prediction error** and build up the standard evaluation toolkit on top of the model we just reproduced.

```
# Analysis frame: same outcome + predictors, complete cases only.
d <- data[, c("TrustCent", predictors)]
d <- d[complete.cases(d), ]
n <- nrow(d)
rmse <- function(a, b) sqrt(mean((a - b)^2))
mae <- function(a, b) mean(abs(a - b))
r2 <- function(truth, pred) 1 - sum((truth - pred)^2) / sum((truth - mean(truth))^2)
cat("Analysis N =", n, "| outcome mean =", round(mean(d$TrustCent), 3),
    "| outcome SD =", round(sd(d$TrustCent), 3), "\n")
```

```
## Analysis N = 1208 | outcome mean = 7.55 | outcome SD = 2.386
```

4.1 Predictive vs. explanatory framing — and why we hold out data

The reproduced regression is fit on all 1208 respondents and then read off its coefficients. If we instead ask “how accurately does this model predict trust?” and answer using the *same* rows we fit on, we flatter ourselves: the model has already seen every answer. The honest question is how it does on respondents it has **never seen**. That is the entire motivation for holding out data. We first show the size of the self-flattery, the **in-sample optimism**.

```
m_full <- lm(form_cent, d)
insamp <- rmse(d$TrustCent, predict(m_full)) # scored on training rows
insampR2 <- summary(m_full)$r.squared
cat("In-sample RMSE (fit and scored on all", n, "rows):", round(insamp, 3), "\n")
```

```
## In-sample RMSE (fit and scored on all 1208 rows): 2.275
```

```
cat("In-sample R^2:", round(insampR2, 3), "\n")
```

```
## In-sample R^2: 0.09
```

That in-sample number is optimistic *by construction*. To measure it honestly we need data the model did not fit. Two protocols follow: a single train/test split (simple, but noisy), then k -fold cross-validation (uses every row for both roles).

4.2 A single train/test split

We randomly assign 70% of respondents to *train* and 30% to *test*, fit OLS on the training rows only, and score on the test rows.

```
tr_idx <- sample(n, round(0.70 * n))
train  <- d[tr_idx, ]
test   <- d[-tr_idx, ]

m_split <- lm(form_cent, train)
tr_rmse <- rmse(train$TrustCent, predict(m_split))      # in-sample (train)
te_rmse <- rmse(test$TrustCent, predict(m_split, test)) # honest (test)

split_tab <- data.frame(
  Quantity = c("Training rows", "Test rows",
               "In-sample RMSE (train)", "Held-out RMSE (test)"),
  Value    = c(nrow(train), nrow(test),
               round(tr_rmse, 3), round(te_rmse, 3)))
knitr::kable(split_tab, caption = "Single 70/30 train/test split, OLS. The held-out RMSE
→ is the honest estimate of prediction error.")
```

Table 12: Single 70/30 train/test split, OLS. The held-out RMSE is the honest estimate of prediction error.

Quantity	Value
Training rows	846.000
Test rows	362.000
In-sample RMSE (train)	2.298
Held-out RMSE (test)	2.228

The in-sample RMSE understates error; the **held-out RMSE is the number to trust**. But a single split wastes 30% of the data and its estimate wobbles depending on which respondents happened to land in the test set. Cross-validation fixes both.

4.3 k -fold cross-validation (folds written out explicitly)

We partition the 1208 respondents into $K = 10$ disjoint folds. Each fold takes a turn as the held-out test set while the other nine train the model; every respondent is predicted exactly once, from a model that never saw them. We print the fold sizes to make the mechanics concrete.

```
K <- 10
folds <- sample(rep(1:K, length.out = n)) # fold label for every respondent
knitr::kable(as.data.frame(table(Fold = folds), responseName = "Size"),
  caption = "Explicit fold assignment: each of the 1208 respondents gets one
→ fold label (1-10).")
```

Table 13: Explicit fold assignment: each of the 1208 respondents gets one fold label (1-10).

Fold	Size
1	121
2	121
3	121
4	121
5	121
6	121
7	121
8	121
9	120
10	120

```
cv_ols_pred <- numeric(n)                                # out-of-fold prediction for each row
for (k in 1:K) {
  tri <- folds != k; tei <- folds == k
  fit <- lm(form_cent, d[tri, ])
  cv_ols_pred[tei] <- predict(fit, d[tei, ])
}
cv_rmse_ols <- rmse(d$TrustCent, cv_ols_pred)
cat("10-fold CV RMSE (OLS):", round(cv_rmse_ols, 3), "\n")
```

```
## 10-fold CV RMSE (OLS): 2.292
```

We now have three numbers for the *same* OLS model: an over-optimistic in-sample RMSE of 2.275, a single-split held-out RMSE of 2.228, and a stable 10-fold CV RMSE of 2.292. The gap between the first and the last is the **optimism** we would have paid for by trusting in-sample fit.

4.4 The bias–variance tradeoff: a validation curve

Prediction error decomposes into **bias** (systematic error from a model too rigid to capture the signal) and **variance** (error from a model so flexible it chases noise in the particular training sample). We expose the tradeoff with a **flexibility knob** that stays entirely within Day-1’s linear-model toolkit: the **degree of a polynomial**. We keep OLS and the same six predictors, but replace the linear income term with a degree- p polynomial in income, `poly(Income2, p)` — income being the only appreciably continuous predictor. Degree $p = 1$ is the original linear fit (most rigid); higher p lets the income–trust curve wiggle more (more flexible). We sweep p and, at each value, record **training** error and **cross-validated** error.

```
degrees <- 1:8
# Build the OLS formula with poly(Income2, p) in place of the linear Income2 term.
poly_form <- function(p) {
  other <- setdiff(predictors, "Income2")
  as.formula(paste("TrustCent ~", paste(other, collapse = " + "),
    "+ poly(Income2, ", p, ")"))
}
poly_train_rmse <- function(p) {
  fit <- lm(poly_form(p), d)                                # fit and scored on all rows (in-sample)
  rmse(d$TrustCent, predict(fit))
}
poly_cv_rmse <- function(p) {
  err <- numeric(K)
  for (f in 1:K) {
```

```

tri <- folds != f; tei <- folds == f
fit <- lm(poly_form(p), d[tri, ])
err[f] <- rmse(d$TrustCent[tei], predict(fit, d[tei, ]))
}
mean(err)
}
vc <- data.frame(degree = degrees,
                 Train = sapply(degrees, poly_train_rmse),
                 CV    = sapply(degrees, poly_cv_rmse))
best_deg <- vc$degree[which.min(vc$CV)]

vc_long <- rbind(data.frame(degree = vc$degree, RMSE = vc$Train, Curve = "Training
  ↪ (in-sample)"),
                 data.frame(degree = vc$degree, RMSE = vc$CV,    Curve = "Cross-validated
  ↪ (held-out)"))
ggplot(vc_long, aes(x = degree, y = RMSE, colour = Curve)) +
  geom_line(linewidth = 0.8) + geom_point() +
  geom_vline(xintercept = best_deg, linetype = "dashed", colour = "grey40") +
  scale_x_continuous(breaks = degrees) +
  labs(x = "polynomial degree in income --- more flexible to the right",
       y = "RMSE", colour = NULL,
       title = paste0("Bias-variance: CV error minimized at degree = ", best_deg)) +
  theme_minimal(base_size = 11) +
  theme(legend.position = "top")

```

Bias-variance: CV error minimized at degree = 3




Figure 5: Validation curve for OLS with a degree- p polynomial in income. Training error falls monotonically with degree (more flexibility); CV (held-out) error is U-shaped, minimized at an intermediate degree. The gap between the curves is the variance the flexible fits pay.

Reading the curve left to right: at degree $p = 1$ the income term is a straight line (most bias, least variance); as the degree grows the training error **falls monotonically** (a more flexible curve always fits the training rows at least as well) while the CV error traces a **U-shape** — **minimized at degree 3** — and then rises as

the high-degree polynomial starts chasing noise (variance). This is the canonical bias–variance tradeoff, and the CV minimum is the flexibility sweet spot.

4.5 Performance metrics for a continuous outcome, vs. a naive baseline

Trust is continuous (0–10), so the relevant metrics are **RMSE** (root mean squared error, in trust points), **MAE** (mean absolute error, in trust points), and R^2 (share of variance explained). Every metric is meaningless without a reference: we always compare against the **naive baseline** that ignores all predictors and simply predicts the training-set mean trust for everyone. All numbers below are 10-fold CV (out-of-fold), so they are honest.

```
# Out-of-fold predictions for baseline, OLS, and the CV-best polynomial model.
base_pred <- numeric(n); poly_pred <- numeric(n)
for (k in 1:K) {
  tri <- folds != k; tei <- folds == k
  base_pred[tei] <- mean(d$TrustCent[tri]) # mean baseline
  fit_p <- lm(poly_form(best_deg), d[tri, ])
  poly_pred[tei] <- predict(fit_p, d[tei, ])
}
metric_tab <- data.frame(
  Model = c("Naive baseline (mean)", "OLS (reproduced spec)",
            paste0("Poly OLS (income degree = ", best_deg, ")")),
  RMSE = round(c(rmse(d$TrustCent, base_pred), cv_rmse_ols, rmse(d$TrustCent,
    ↪ poly_pred)), 3),
  MAE = round(c(mae(d$TrustCent, base_pred),
                mae(d$TrustCent, cv_ols_pred),
                mae(d$TrustCent, poly_pred)), 3),
  R2 = round(c(r2(d$TrustCent, base_pred),
                r2(d$TrustCent, cv_ols_pred),
                r2(d$TrustCent, poly_pred)), 3))
knitr::kable(metric_tab, caption = "10-fold CV metric suite vs. a mean baseline. RMSE and
↪ MAE in trust points (0-10 scale); R^2 is out-of-sample.")
```

Table 14: 10-fold CV metric suite vs. a mean baseline. RMSE and MAE in trust points (0-10 scale); R^2 is out-of-sample.

Model	RMSE	MAE	R2
Naive baseline (mean)	2.387	1.895	-0.002
OLS (reproduced spec)	2.292	1.797	0.077
Poly OLS (income degree = 3)	2.283	1.789	0.084

The baseline’s out-of-sample R^2 is ~ 0 by construction (predicting the mean explains no variance beyond the mean). The reproduced OLS and the CV-best polynomial both beat the baseline on all three metrics, but only modestly — and the polynomial barely improves on the plain linear fit, a preview of the closing point that six mostly-categorical predictors carry a **small, largely linear** signal.

4.6 Parameters vs. hyperparameters; grid and random search

A **parameter** is learned by fitting on data — the OLS coefficients $\hat{\beta}$, learned by least squares. A **hyperparameter** is a knob *we* set before fitting and cannot read off the data by least squares: here it is the **polynomial degree**, and we can also treat **which predictor** receives the polynomial as a second knob. We tune hyperparameters by scoring each candidate configuration with cross-validation. Two standard search strategies:


```

# CV RMSE for a polynomial configuration: (degree p, target predictor), fixed folds.
poly_form_var <- function(p, var) {
  other <- setdiff(predictors, var)
  as.formula(paste("TrustCent ~", paste(other, collapse = " + "),
    "+ poly(", var, ",", p, ")"))
}
cv_poly_cfg <- function(p, var) {
  err <- numeric(K)
  for (f in 1:K) {
    tri <- folds != f; tei <- folds == f
    fit <- lm(poly_form_var(p, var), d[tri, ])
    err[f] <- rmse(d$TrustCent[tei], predict(fit, d[tei, ]))
  }
  mean(err)
}
# Max usable degree per predictor: a degree-p polynomial needs > p distinct values.
max_deg <- function(var) length(unique(d[[var]])) - 1

```

Grid search evaluates every combination on a predefined lattice. We cross polynomial degree with the two predictors that have enough distinct values to support one (Income2, up to degree 8; Education, which has only 3 levels, up to degree 2).

```

grid_var <- c("Income2", "Education")
grid <- do.call(rbind, lapply(grid_var, function(v)
  data.frame(degree = 1:min(8, max_deg(v)), var = v, stringsAsFactors = FALSE)))
grid$cvRMSE <- mapply(cv_poly_cfg, grid$degree, grid$var)
grid_best <- grid[which.min(grid$cvRMSE), ]
cat("Grid search evaluated", nrow(grid), "configurations. Best:",
  "degree =", grid_best$degree, "| predictor =", grid_best$var,
  "| CV RMSE =", round(grid_best$cvRMSE, 3), "\n")

```

```

## Grid search evaluated 10 configurations. Best: degree = 3 | predictor = Income2 | CV
↪ RMSE = 2.275

```

Random search samples configurations at random from the same ranges — often finding a comparable optimum with fewer evaluations, and here randomizing both the degree and which predictor gets the polynomial.

```

set.seed(7)
n_rand <- 15
rand <- data.frame(var = sample(grid_var, n_rand, replace = TRUE),
  stringsAsFactors = FALSE)
# Draw a valid degree for each sampled predictor (respecting its distinct-value cap).
rand$degree <- vapply(rand$var, function(v) sample(1:min(8, max_deg(v)), 1), integer(1))
rand$cvRMSE <- mapply(cv_poly_cfg, rand$degree, rand$var)
rand_best <- rand[which.min(rand$cvRMSE), ]
cat("Random search evaluated", n_rand, "configurations. Best:",
  "degree =", rand_best$degree, "| predictor =", rand_best$var,
  "| CV RMSE =", round(rand_best$cvRMSE, 3), "\n")

```

```

## Random search evaluated 15 configurations. Best: degree = 3 | predictor = Income2 | CV
↪ RMSE = 2.275

```

Validation curves from the two searches — one line per predictor (grid) plus the random draws overlaid — make the tuning visible:

```
ggplot() +
  geom_line(data = grid, aes(x = degree, y = cvRMSE, colour = var), linewidth = 0.8) +
  geom_point(data = grid, aes(x = degree, y = cvRMSE, colour = var)) +
  geom_point(data = rand, aes(x = degree, y = cvRMSE, shape = "random draw"),
    colour = "black", size = 2) +
  geom_point(data = grid_best, aes(x = degree, y = cvRMSE), colour = "red",
    size = 4, shape = 1, stroke = 1.2) +
  scale_shape_manual(values = c("random draw" = 4)) +
  scale_x_continuous(breaks = 1:8) +
  labs(x = "polynomial degree", y = "CV RMSE", colour = "predictor (grid)", shape = NULL,
    title = "Grid (lines) and random (crosses) search; red ring = grid optimum") +
  theme_minimal(base_size = 11) + theme(legend.position = "top")
```

Grid (lines) and random (crosses) search; red ring = grid optimum




Figure 6: Validation curves from the hyperparameter searches. Grid: CV RMSE vs. polynomial degree for each candidate predictor (lines). Random: sampled configurations (crosses). Both searches converge on the same shallow optimum.

Both strategies land on essentially the same shallow optimum (CV RMSE around 2.275), confirming that the tuning surface is flat — there is no exotic polynomial degree that unlocks hidden predictive power here.

4.7 Tying evaluation back to the reproduced regression

Finally, we put the honestly-evaluated learners side by side with the reproduced OLS and the naive baseline, all under the identical 10-fold protocol.

```
# Out-of-fold predictions for the CV-best polynomial from the grid search.
poly_best_pred <- numeric(n)
for (k in 1:K) {
  tri <- folds != k; tei <- folds == k
  fit_b <- lm(poly_form_var(grid_best$degree, grid_best$var), d[tri, ])
  poly_best_pred[tei] <- predict(fit_b, d[tei, ])
}
```

```

final_tab <- data.frame(
  Model = c("Naive baseline (mean)",
    "OLS --- reproduced Table-5 spec",
    paste0("Poly OLS --- best grid (", grid_best$var, ", degree ",
      ↪ grid_best$degree, ")")),
  CV_RMSE = round(c(rmse(d$TrustCent, base_pred),
    cv_rmse_ols,
    grid_best$cvRMSE), 3),
  CV_R2 = round(c(r2(d$TrustCent, base_pred),
    r2(d$TrustCent, cv_ols_pred),
    r2(d$TrustCent, poly_best_pred)), 3))
knitr::kable(final_tab, caption = "Honest (10-fold CV) comparison: the reproduced OLS,
↪ the CV-best polynomial model, and the naive baseline predicting central-government
↪ trust.")

```

Table 15: Honest (10-fold CV) comparison: the reproduced OLS, the CV-best polynomial model, and the naive baseline predicting central-government trust.

Model	CV_RMSE	CV_R2
Naive baseline (mean)	2.387	-0.002
OLS — reproduced Table-5 spec	2.292	0.077
Poly OLS — best grid (Income2, degree 3)	2.275	0.084

The reproduced regression — the very model whose -0.511 Tiananmen-cohort coefficient is the paper’s headline — is **not just an explanatory device; it is also a competitive predictor**. After honest cross-validation, OLS matches the CV-best polynomial model and both clear the mean baseline by a similar, modest margin. The predictive lens does not overturn the explanatory story: it shows that the signal the paper documents (a Tiananmen-cohort gap of about half a point, plus the education/CCP/income controls) is close to *all the predictable signal there is* in these six variables. The gap between the in-sample RMSE (2.275) and the CV RMSE (2.292) is small precisely because a six-variable, near-linear model has little room to overfit — which is exactly why the paper’s OLS inference is trustworthy.

5 Takeaway

Day 1’s toolkit, demonstrated end to end on the reproduced trust model:

- **Explanatory vs. predictive framing.** The paper estimates and interprets $\beta_T = -0.511$; we re-asked how well the same model *predicts unseen respondents*.
- **In-sample optimism.** In-sample RMSE 2.275 vs. honest CV RMSE 2.292 — the optimism gap made explicit.
- **Data splitting.** A single 70/30 split, then explicit 10-fold CV with the fold sizes printed.
- **Bias–variance.** A polynomial-degree validation curve: training error falls monotonically with flexibility, CV error U-shaped, minimized at degree 3.
- **Metric suite + baseline.** RMSE, MAE and R^2 for the continuous outcome, always against the mean baseline.
- **Hyperparameters vs. parameters, grid & random search.** Polynomial degree (and which predictor gets the polynomial) tuned by both grid and random search, with validation curves from each.
- **Back to the regression.** Under honest CV the reproduced OLS ties the CV-best polynomial model and beats the baseline: the paper’s explanatory model is also a sound predictive one, and its low overfitting risk is *why* its inference holds up.

6 Independent exercises (each about 10 minutes)

1. **How much work does one control do?** Refit the reproduced specification *without* `Education`, reusing the same `folds` (copy the loop from the final-comparison chunk and edit the formula). Compare its out-of-fold RMSE with `cv_rmse_ols` and the mean baseline in `final_tab`.
2. **The headline variable alone.** Add a `TrustCent ~ T` row to the same comparison. How much of the six-variable model's edge over the baseline does the Tiananmen-cohort indicator deliver by itself?
3. **CV design.** Set `K <- 5`, redraw `folds`, and recompute `cv_rmse_ols`. How much does the CV estimate move, and why would you expect 5-fold and 10-fold CV to disagree only slightly on `n` of this size?